


PREVIOUSLY ON <sup>1</sup><sub>^</sub>24 & 224

# CLIENT AND SERVER SOCKETS (SYSTEM CALLS)



# SENDING DATA (BOTH CLIENT AND SERVER)

```
char *data_addr = "hello, world";
int data_len = 12;

int sent_bytes = send(sock, data_addr, data_len, 0);
if (sent_bytes < 0) {
    perror("send failed");
}
```

- **send()**: sends data, returns the number of sent bytes
  - Also OK with `write()`, `writew()`
- **data\_addr**: address of data to send
- **data\_len**: size of the data
- With blocking sockets (default), `send()` blocks until it sends all the data.
- With non-blocking sockets, **sent\_bytes** may not equal to **data\_len**
  - If kernel does not have enough space, it accepts only partial data
  - You must retry for the unsent data

# RECEIVING DATA (BOTH CLIENT AND SERVER)

```
char buffer[4096];
int expected_data_len = sizeof(buffer);

int read_bytes = recv(sock, buffer, expected_data_len, 0);
if (read_bytes == 0) { // connection is closed
    ...
} else if (read_bytes < 0) { // error
    perror("recv failed");
} else { // OK. But no guarantee read_bytes == expected_data_len
    ...
}
```


- **recv()**: reads bytes from the socket and returns the number of read bytes.
- **read\_bytes** may not equal to **expected\_data\_len**
  - If no data is available, it blocks
  - If only partial data is available, `read_bytes < expected_data_len`
  - On socket close, `expected_data_len` equals to 0 (not error!)
  - If you get only partial data, you should retry for the remaining portion.
- **TIP: You almost always want to call `recv()/Read()` inside a loop...**



# Today's agenda

- Let's write a couple Go programs that send/receive data over the network using the TCP protocol
- Then we'll go back and take a deep dive on why they work and what is going on "under the engine hood"
- TCP and "blocking" in the OS


# CONCEPTUAL SETUP



# DIGGING INTO SEND() A BIT MORE


```
rv = connect(s,...);  
:  
:  
rv = send(s,buffer0,1000,0);  
:  
:  
rv = send(s,buffer1,2000,0);  
:  
:  
rv = send(s,buffer2,5000,0);  
:  
:  
close(s);
```

# AFTER 3 SEND() CALLS






# AFTER FIRST RECV()



# AFTER ANOTHER RECV()

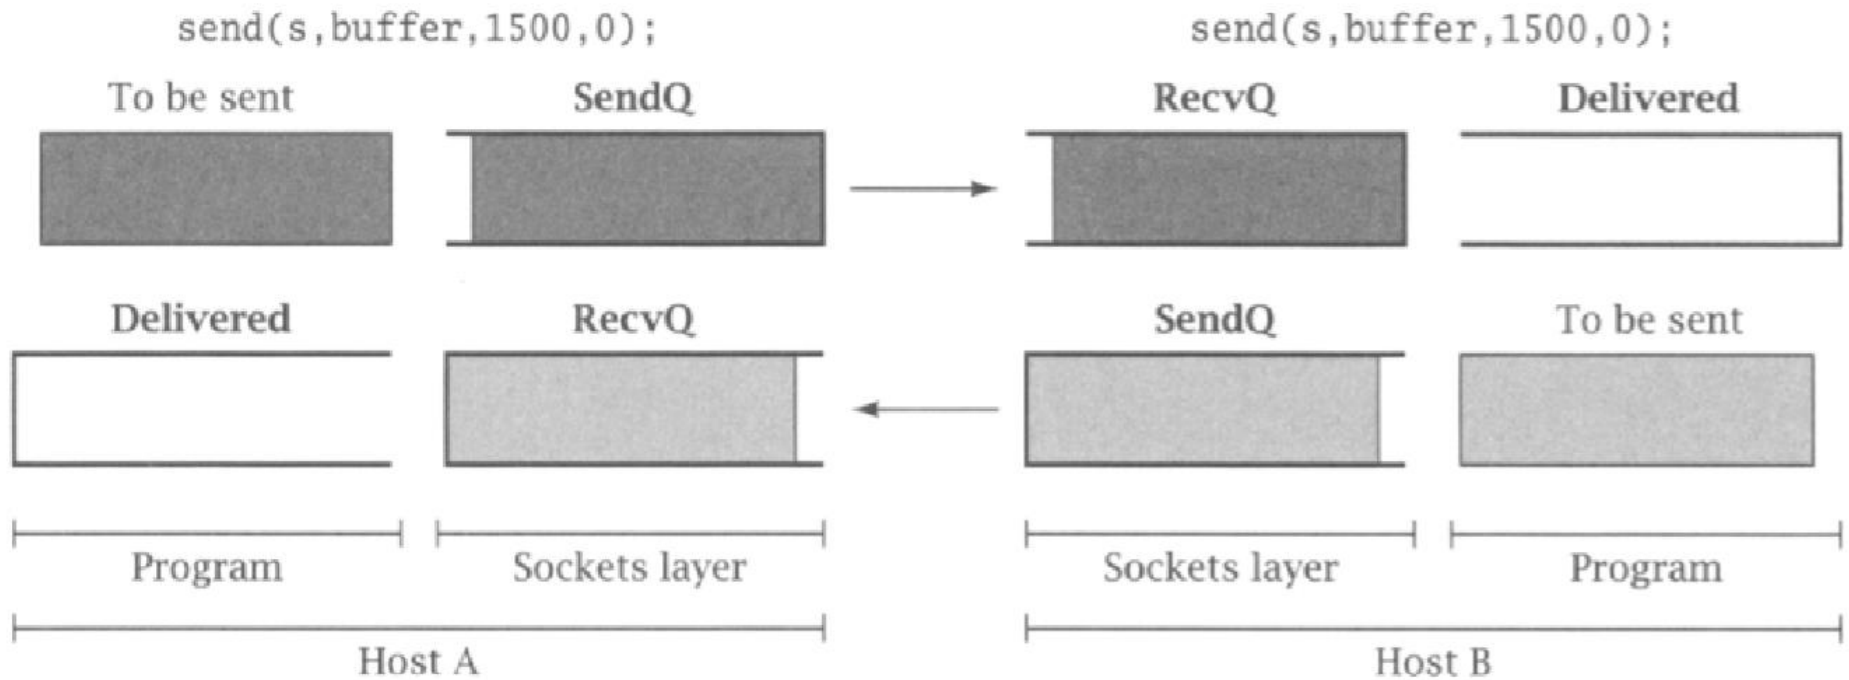


- 1** First send call (1000 bytes)
- 2** Second send call (2000 bytes)
- 3** Third send call (5000 bytes)

# WHEN DOES BLOCKING OCCUR?

- SendQ size: **SQS**
- RecvQ size: **RQS**
- `send(s, buffer, n, 0);`
  - $n > \text{SQS}$ : blocks until  $(n - \text{SQS})$  bytes xfered to RecvQ
  - If  $n > (\text{SQS} + \text{RQS})$ , blocks until receiver calls `recv()` enough to read in  $n - (\text{SQS} + \text{RQS})$  bytes
- How does this lead to deadlock?
  - Trivial cause: both sides call `recv()` w/o sending data

# MORE SUBTLE REASON FOR DEADLOCK



- SendQ size = 500; RecvQ size = 500

**CSE 124 AND CSE 224:**

**SENDING AND RECEIVING DATA, ERROR  
HANDLING, AND PRACTICAL  
CONSIDERATIONS**

George Porter  
April 15, 2025



# ATTRIBUTION

- These slides are released under an Attribution-NonCommercial-ShareAlike 3.0 Unported (CC BY-NC-SA 3.0) Creative Commons license
- These slides incorporate material from:
  - Computer Networks: A Systems Approach, 5e, by Peterson and Davie
  - Michael Freedman and Kyle Jamieson, Princeton University (also under a CC BY-NC-SA 3.0 Creative Commons license)

# WHAT ARE PROTOCOLS?

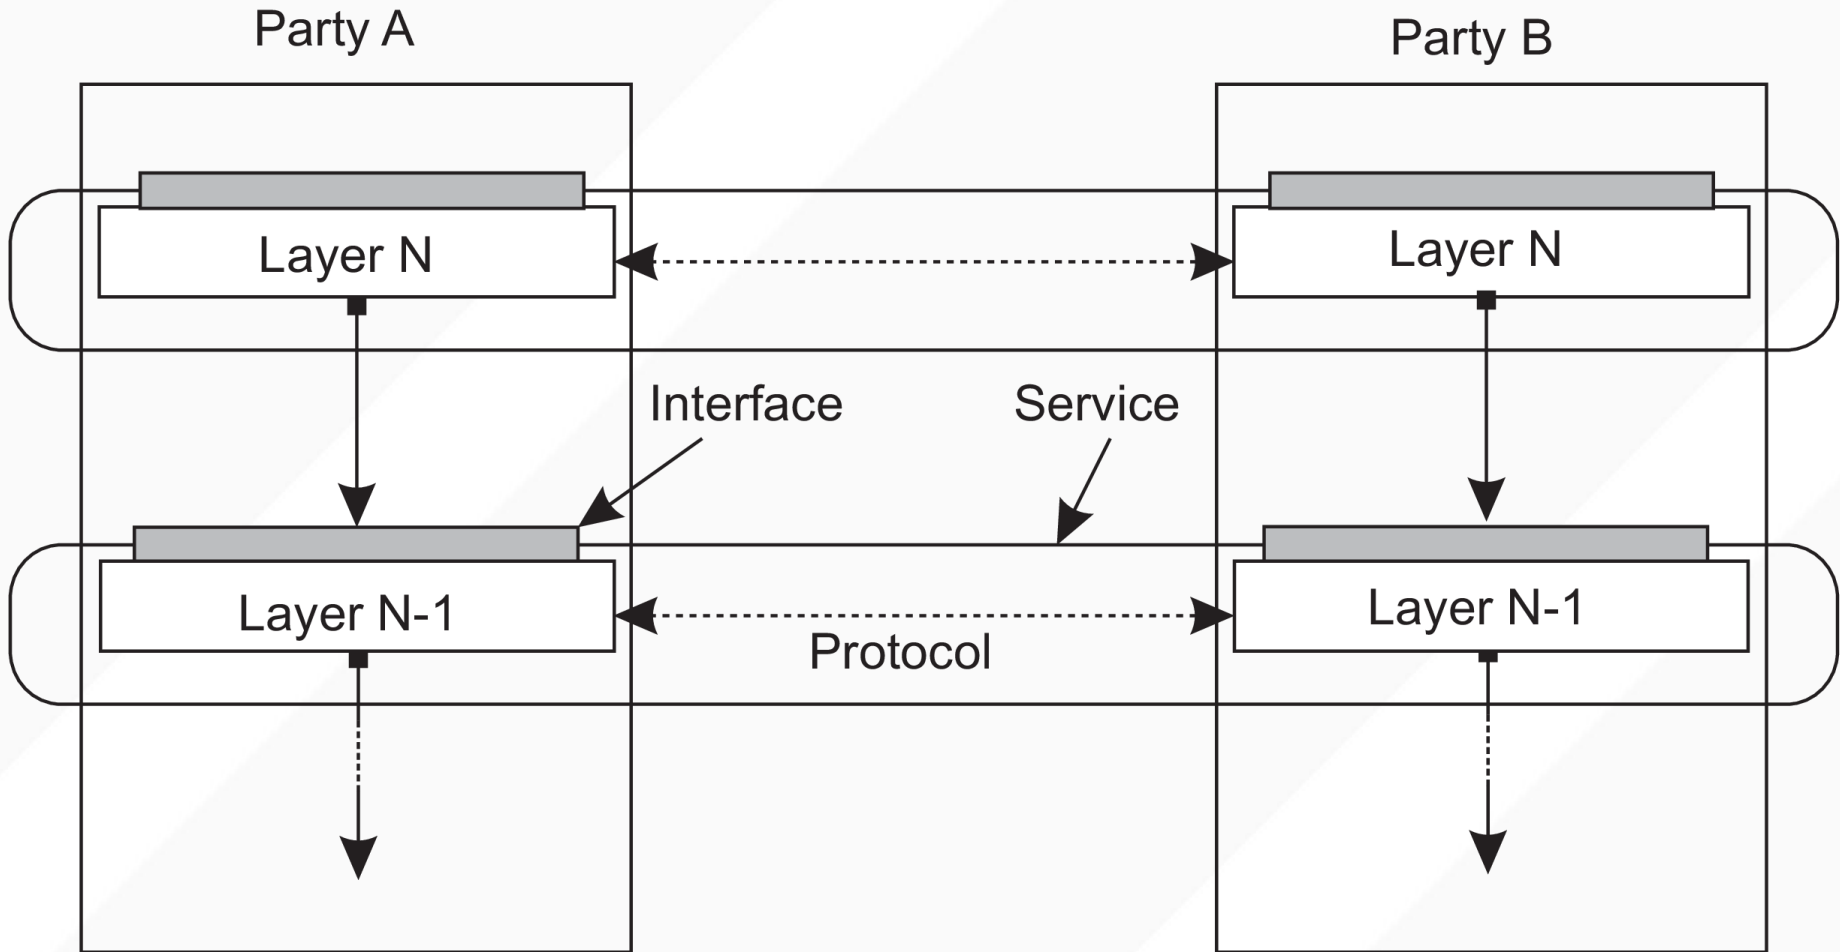


Image from public domain

- Explicit and implicit conventions for how to communicate
  - Not for what is communicated
- Enables heterogeneous architectures, languages, OSes, byte ordering, ...



# SERVICE AND PROTOCOL INTERFACES





# WHERE DO PROTOCOLS COME FROM?

- Standards bodies
  - IETF: Internet Engineering Task Force
  - ISO: International Standards Organization
- Community efforts
  - “Request for comments”
  - Bitcoin
- Corporations/industry
  - RealAudio™, Call of Duty multiplayer, Skype



International  
Organization for  
Standardization



# HOW ARE PROTOCOLS SPECIFIED?

## Prose/BNF

### 3.2. HEADER FIELD DEFINITIONS

These rules show a field meta-syntax, without regard for the particular type or internal syntax. Their purpose is to permit detection of fields; also, they present to higher-level parsers an image of each field as fitting on one line.

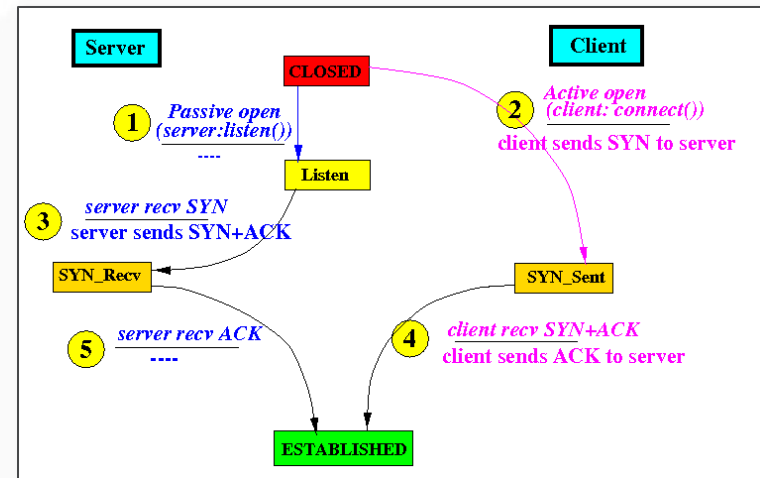
field = field-name ":" [ field-body ] CRLF

field-name = 1\*<any CHAR, excluding CTLs, SPACE, and ":">


field-body = field-body-contents  
[CRLF LWSP-char field-body]

field-body-contents =  
<the ASCII characters making up the field-body, as  
defined in the following sections, and consisting  
of combinations of atom, quoted-string, and  
specials tokens, or else consisting of texts>

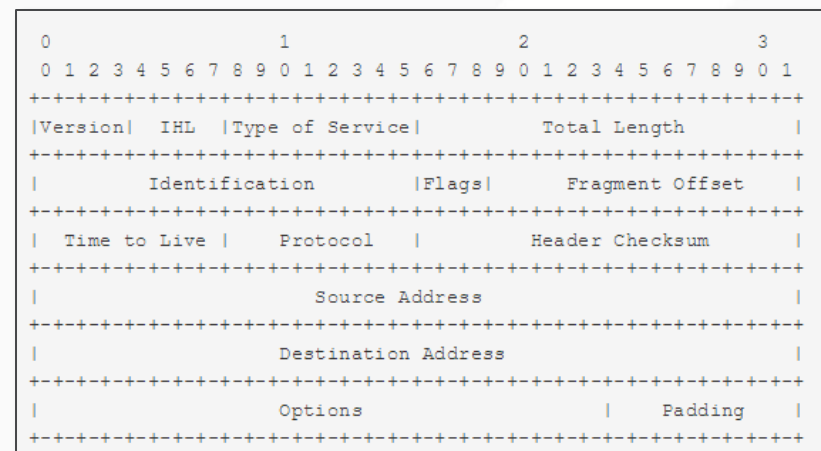
## State transition diagrams



## Message Sequence Diagram



## Packet formats



# EXAMPLE: A SIMPLE VOTE COUNTING SYSTEM


# DEFINITIONS

- Operation (e.g., in a voting system)
  - An action you can perform within a protocol's service interface
  - E.g., "Submit vote", "get current vote count", "reset vote count to zero"
- Message
  - An encoding of an operation or data according to a protocol's *wire format*. Common formats include XML, binary, JSON, ...
- Framing
  - Writing out (and reading in) messages from a stream such that messages can be separated and interpreted correctly
- Parsing/encoding/decoding
  - Converting a message to/from an application-level data structure

# PARSING: CONVERTING IN-MEMORY REPRESENTATION WITH A “WIRE” REPRESENTATION

```
type Employee struct {  
    operation uint8  
    id uint64  
    department uint16  
}
```

- Binary



- Text (ad-hoc)

“OP=1, id=428, d=80”

(1,428,80)

- Text (XML)

```
<employee>  
  <operation>1</operation>  
  <id>428</id>  
  <department>80</department>  
</employee>
```

- Many others...

# FRAMING: LENGTH SPECIFICATION VS DELIMITERS

```
package main

import "fmt"

type Employee struct {
    operation uint8
    name string
    id uint64
    department uint16
}
```

- Binary representation of name?
  - Handling variable length
- Consider “Alan” as a name

97	108	97	110
----	-----	----	-----

- Option 1: Explicit length

4	97	108	97	110
---	----	-----	----	-----

- But how big should length be?

- Option 2: Delimiter

97	108	97	110	0
----	-----	----	-----	---

- But what if delimiter is in the message?

# FRAMING: DETERMINING MESSAGE BOUNDARIES

00009a0	3053	3720	2035	2030	2f52	5343	2031	3537
00009b0	3020	5220	3e3e	452f	7478	5347	6174	6574
00009c0	3c3c	472f	3053	3720	2036	2030	2f52	5347
00009d0	2031	3237	3020	5220	3e3e	462f	6e6f	3c74
00009e0	2f3c	3243	305f	3720	2033	2030	2f52	3154
00009f0	305f	3820	2034	2030	2f52	5454	2030	3437
0000a00	3020	5220	542f	3154	3720	2038	2030	2f52
0000a10	5454	2032	3338	3020	5220	3e3e	502f	6f72
0000a20	5363	7465	2f5b	4450	2f46	6554	7478	492f
0000a30	616d	6567	5d43	502f	6f72	6570	7472	6569
0000a40	3c73	2f3c	434d	2030	3236	3020	5220	3e3e
0000a50	582f	624f	656a	7463	3c3c	492f	306d	3920
0000a60	2039	2030	2f52	6d49	2031	3031	2031	2030
0000a70	3e52	3e3e	2f3e	6f52	6174	6574	3020	542f
0000a80	7568	626d	3520	2032	2030	2f52	7254	6d69
0000a90	6f42	5b78	2e30	2030	2e30	2030	3136	2e32
0000aa0	2030	3937	2e32	5d30	542f	7079	2f65	6150
0000ab0	6567	3e3e	650d	646e	626f	0d6a	3836	3020
0000ac0	6f20	6a62	3c0d	2f3c	6946	746c	7265	462f
0000ad0	616c	6574	6544	6f63	6564	4c2f	6e65	7467
0000ae0	2068	3333	3e32	733e	7274	6165	0d6d	480a

- Framing

- Finds and returns bytes corresponding to single message
- Even if messages are variable length
- Writes out bytes corresponding to a message with enough *context* for the other side to determine the message boundaries

# FRAMING SCENARIO

- Consider a voting scenario
- Each message is variable length
  - “Voting v 134” → [Vote for candidate 134]
  - “Voting i 19381”
    - → [Query candidate 19381’s vote count]
  - First is 12 characters, second is 14 characters
- Given a stream of vote operations, how to separate them?



# FRAMING CHOICES

- **Delimiter (in this case '\$')**

```
Voting v 134$Voting v 2817$Voting i  
9172651$Voting v 2$Voting i 1900$Voting v  
32$Voting i 8
```

- **Length + message**

```
12Voting v 13413Voting v 281716Voting i  
917265110Voting v 213Voting i 190011Voting v  
3210Voting i 8
```

# THE MAIN LOOP OF YOUR SERVER

```
Remaining := ""
```

```
buf := make([]byte, 1024)
```

```
for {
```

```
    for "Does remaining contain a full request?" {
```

```
        If yes, (1) parse it, then (2) remove from remaining
```

```
    }
```

```
    size, err := c.Read(buf)
```

```
    data := buf[:size]
```

```
    remaining = remaining + string(data)
```

```
}
```

*How do you know when a request is completed?*

# HOW TO TELL IF BUFFER CONTAINS A COMPLETE REQUEST?

- This is the framing problem
- For length-based framing:

12	97	108	97	...
----	----	-----	----	-----

- Keep reading until we have 12 bytes of request data
- For delimiter-based framing:
  - OK to simply scan for delimiters using e.g., a for loop

# FRAMING: SUMMARY

- Writing data

- Given an array of bytes representing an application-level operation, writes to stream

1. Explicit length

- Writes out the length of the message, then message

2. Delimiter

- Ensures delimiter doesn't appear in message
- Writes out message
- Then writes out delimiter

- Reading data

- Reads from stream until entire message is read, returns to higher layer

1. Explicit length

- Reads the length, then reads that many bytes (security?)

2. Delimiter

- Reads continuously into a buffer until delimiter is encountered
- Message then returned to higher layer

# PRIMARY FRAMING/PARSING LOOP

- [see [turing-printer.go](#) and [turing-sender.go](#) demo]
- Concepts introduced in this code:
  - flags package
  - defer command
  - pointers
  - 'go' command (for concurrency)
  - Anonymous functions
  - bufio package and Scanner

UC San Diego